
Week 15

Basic Data Structures
(Chapter 16)

Chapter Goals



© andrea laurita/iStockphoto.

- To understand the implementation of linked lists and array lists
- To analyze the efficiency of fundamental operations of lists and arrays
- To implement the stack and queue data types
- To implement a hash table and understand the efficiency of its operations

Implementing Linked Lists - The Node Class

- We will implement a simplified, singly-linked list.
- A linked list stores elements in a sequence of nodes.
- A `Node` object stores an element and a reference to the next node.
 - private inner class
 - public instance variables

```
public class LinkedList
{
    class Node
    {
        public Object data;
        public Node next;
    }
}
```

Implementing Linked Lists - The Node Class

- A linked list object holds a reference to the first node:
 - Each node holds a reference to the next node.

```
public class LinkedList
{
    private Node first;

    public LinkedList() { first = null; }

    public Object getFirst()
    {
        if (first == null) { throw new NoSuchElementException(); }
        return first.data;
    }
}
```

Implementing Linked Lists - Adding and Removing the First Element

- When adding or removing the first element, the reference to the first node must be updated.

```
public class LinkedList
{
    . . .
    public void addFirst(Object element)
    {
        Node newNode = new Node(); ❶
        newNode.data = element;
        newNode.next = first; ❷
        first = newNode; ❸
    }
    . . .
}
```

Implementing Linked Lists - Adding the First Element

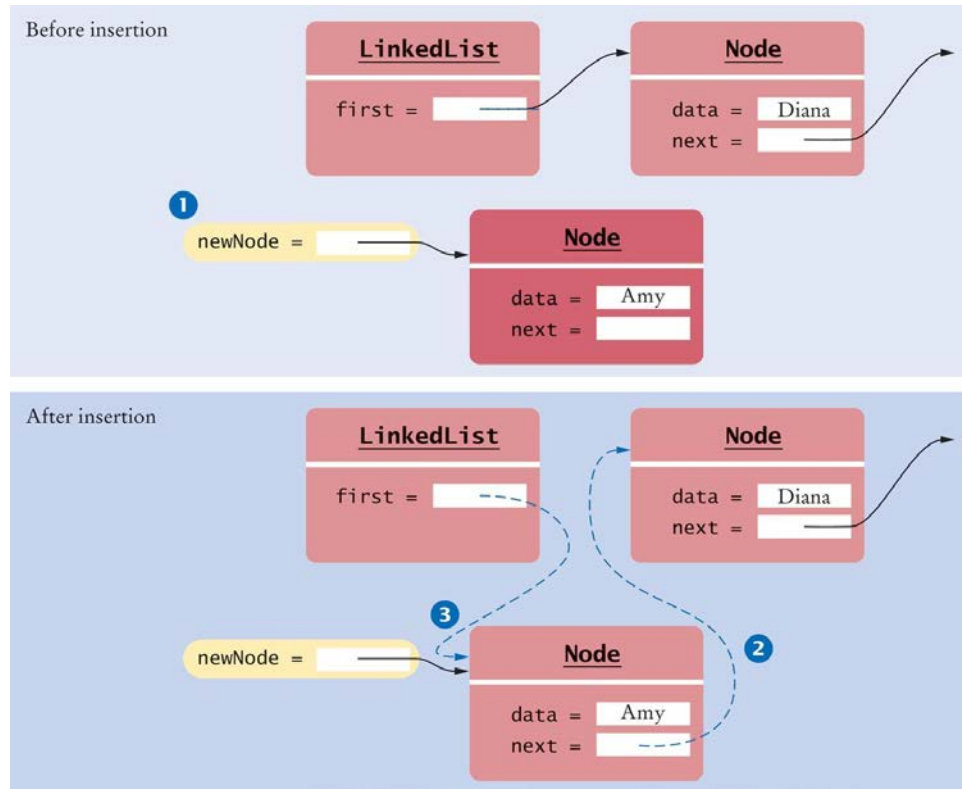


Figure 1 Adding a Node to the Head of a Linked List

Implementing Linked Lists - Removing the First Element

- The data of the first node are saved and later returned as the method result.
- The successor of the first node becomes the first node of the shorter list.
 - The old node is eventually recycled by the garbage collector.

```
public class LinkedList
{
    . . .
    public Object removeFirst()
    {
        if (first == null) { throw new NoSuchElementException(); }
        Object element = first.data;
        first = first.next; ❶
        return element;
    }
    . . .
}
```

Implementing Linked Lists - Removing the First Element

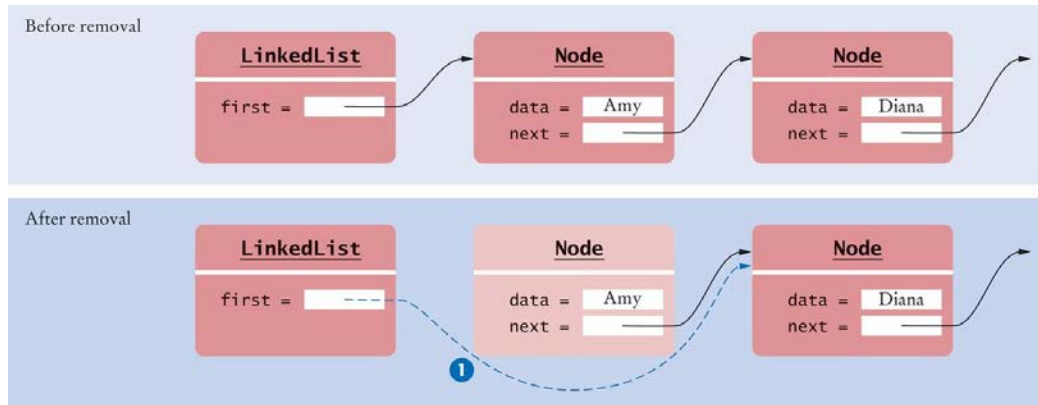


Figure 2 Removing the First Node from a LinkedList

The Iterator Class

- Our simplified `ListIterator` interface has methods: `next`, `hasNext`, `remove`, `add`, and `set`.
- Our `LinkedList` class declares a private inner class `LinkedListIterator`.
 - `LinkedListIterator` implements our simplified `ListIterator` interface.
 - As an inner class `LinkedListIterator` has access to
 - The instance variable `first`
 - The private `Node` class.
- A list iterator object has:
 - A reference to the the currently visited node, `position`
 - A reference to the last node before that, `previous`
 - A `isAfterNext` flag to track when the `next` method has been called.

The Iterator Class

- The LinkedListIterator class:

```
public class LinkedList
{
    . . .
    public ListIterator listIterator()
    {
        return new LinkedListIterator();
    }
    class LinkedListIterator implements ListIterator
    {
        private Node position;
        private Node previous;
        private boolean isAfterNext;
        public LinkedListIterator()
        {
            position = null;
            previous = null;
            isAfterNext = false;
        }
        . . .
    }
}
```

Advancing an Iterator

- To advance an iterator:
 - Update the position
 - Remember the old position for the `remove` method.
- The `next` method:

```
class LinkedListIterator implements ListIterator
{
    . . .
    public Object next()
    {
        if (!hasNext()) { throw new NoSuchElementException(); }
        previous = position; // Remember for remove
        isAfterNext = true;

        if (position == null)
        { position = first; }
        else
        { position = position.next; }
        return position.data;
    }
    . . .
}
```

Advancing an Iterator

- The iterator is at the end
 - if the list is empty (`first == null`) or
 - if there is no element after the current position (`position.next == null`).
- The `hasNext` method:

```
class LinkedListIterator implements ListIterator
{
    . . .
    public boolean hasNext()
    {
        if (position == null)
        {
            return first != null;
        }
        else
        {
            return position.next != null;
        }
    }
    . . .
}
```

Removing an Element

- If this is the first element:
 - Call `removeFirst`
 - Otherwise, update the `next` reference of the previous node
- Update `isAfterNext` to disallow another call to `remove`.
- The `remove` method:

```
class LinkedListIterator implements ListIterator
{
    . . .
    public void remove()
    {
        if (!isAfterNext) { throw new IllegalStateException(); }
        if (position == first)
        {
            removeFirst();
        }
        else
        {
            previous.next = position.next; ❶
        }
        position = previous; ❷
        isAfterNext = false; ❸
    }
    . . .
}
```

Removing an Element

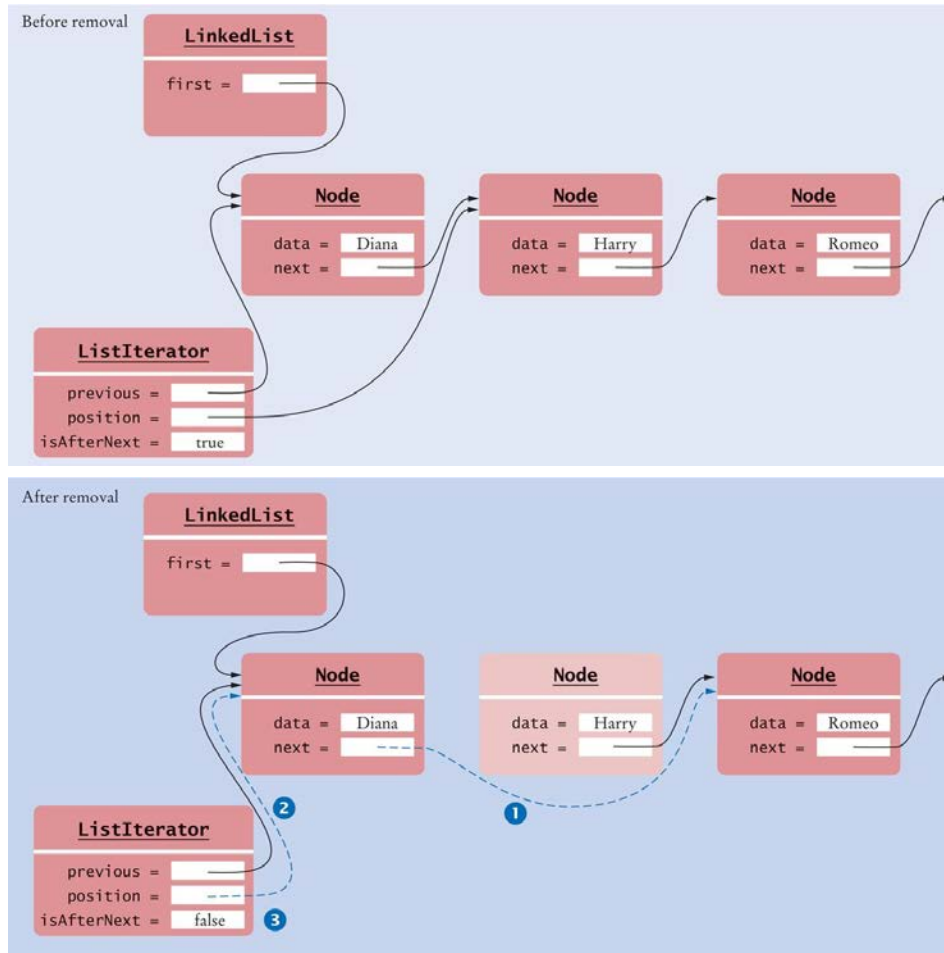


Figure 3 Removing a Node from the Middle of a Linked List

Adding an Element

- After adding the new element
 - set the `isAfterNext` flag to false to disallow a subsequent call to the `remove` or `set` method
- The `add` method:

```
class LinkedListIterator implements ListIterator
{
    . . .
    public void add(Object element)
    {
        if (position == null)
        {
            addFirst(element);
            position = first;
        }
        else
        {
            Node newNode = new Node();
            newNode.data = element;
            newNode.next = position.next; ①
            position.next = newNode; ②
            position = newNode; ③
        }
        isAfterNext = false; ④
    }
    . . .
}
```

Adding an Element

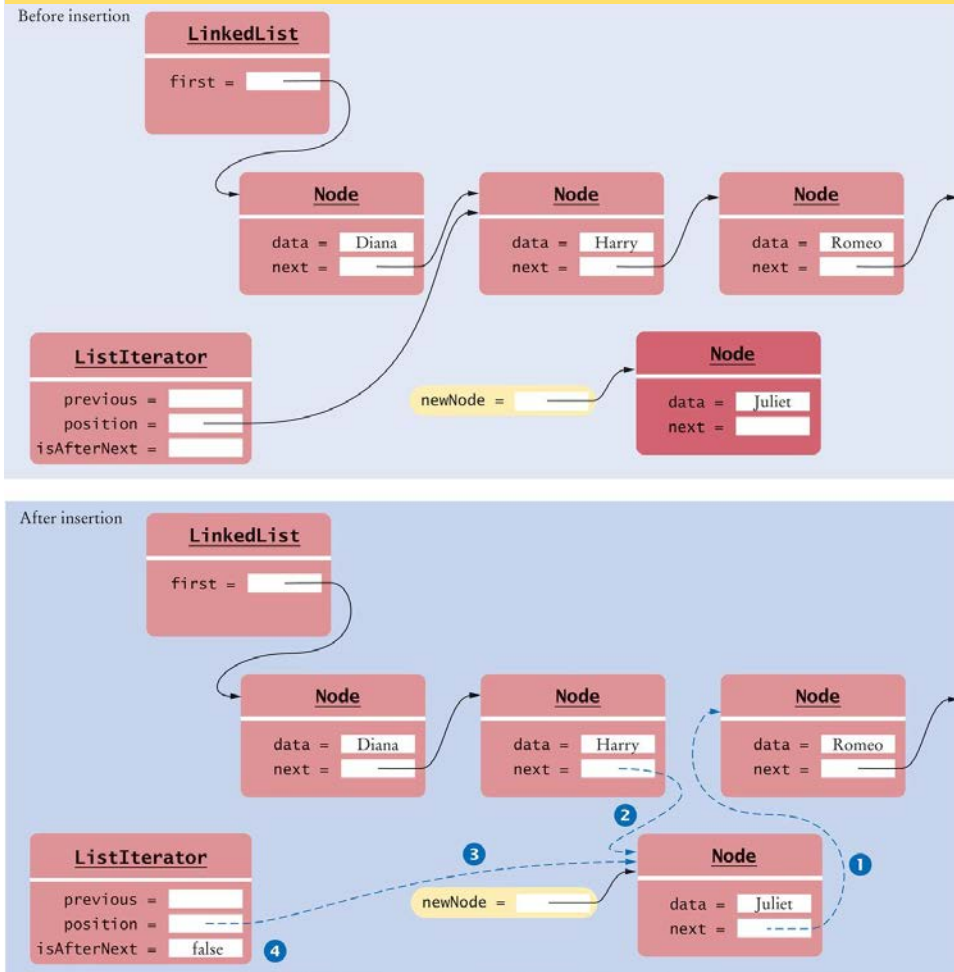


Figure 4 Adding a Node to the Middle of a Linked List

Setting an Element to a Different Value

- `set` method changes the data in the previously visited element.
- Must follow a call to `next`.
- The `set` method:

```
public void set(Object element)
{
    if (!isAfterNext) { throw new IllegalStateException(); }
    position.data = element;
}
```

Efficiency of Linked List Operations

- To get the k^{th} element of a linked list, you start at the beginning of the list and advance the iterator k times.
- To get to the k^{th} node of a linked list, one must skip over the preceding nodes.



© Kris Hanke/iStockphoto.

Efficiency of Linked List Operations

- When adding or removing an element, we update a couple of references in a constant number of steps.
- **Adding and removing an element at the iterator position in a linked list takes $O(1)$ time.**

Efficiency of Linked List Operations

- To add an element at the end of the list
 - Must get to the end - an $O(n)$ operation
 - Add the element $O(1)$ operation
- Adding to the end of a linked list in our implementation takes $O(n)$ time
If the linked list keeps a reference to `last` as well as `first`
 - The time is reduced to constant time: $O(1)$

```
public class LinkedList
{
    private Node first;
    private Node last;
    . . .
}
```

- **We will conclude that adding to the end of a linked list is $O(1)$.**

Efficiency of Linked List Operations

- To remove an element from the end of the list:
 - Need a reference to the next-to-last element so that we can set its `next` reference to null
 - Takes $n-1$ iterations
- Removing an element from the end of the list is $O(n)$.

Efficiency of Linked List Operations




Figure 5 Removing the Last Element of a Singly-Linked List

Efficiency of Linked List Operations

- In a doubly-linked list, each node has a reference to the previous node in addition to the next one.

```
public class LinkedList
{
    . . .
    class Node
    {
        public Object data;
        public Node next;
        public Node previous;
    }
}
```

Efficiency of Linked List Operations

- In a doubly-linked list, removal of the last element takes a constant number of steps.

```
last = last.previous; ❶  
last.next = null;    ❷
```


Efficiency of Linked List Operations




Figure 6 Removing the Last Element of a Doubly-Linked List

Efficiency of Linked List Operations

Table 1 Efficiency of Linked List Operations

Operation	Singly-Linked List	Doubly-Linked List
Access an element.	$O(n)$	$O(n)$
Add/remove at an iterator position.	$O(1)$	$O(1)$
Add/remove first element.	$O(1)$	$O(1)$
Add last element.	$O(1)$	$O(1)$
Remove last element.	$O(n)$	$O(1)$

section_1/LinkedList.java

```
1  import java.util.NoSuchElementException;
2
3  /**
4   A linked list is a sequence of nodes with efficient
5   element insertion and removal. This class
6   contains a subset of the methods of the standard
7   java.util.LinkedList class.
8  */
9  public class LinkedList
10 {
11     private Node first;
12
13     /**
14      Constructs an empty linked list.
15     */
16     public LinkedList()
17     {
18         first = null;
19     }
20
21     /**
22      Returns the first element in the linked list.
23      @return the first element in the linked list
24     */
25     public Object getFirst()
26     {
27         if (first == null) { throw new NoSuchElementException();
28         } return first.data;
29     }
30
31     /**
32      Removes the first element in the linked list.
33      @return the removed element
34     */
35 }
```

section_1/ListIterator.java

```
1  /**
2   A list iterator allows access of a position in a linked list.
3   This interface contains a subset of the methods of the
4   standard java.util.ListIterator interface. The methods for
5   backward traversal are not included.
6  */
7  public interface ListIterator
8  {
9      /**
10       Moves the iterator past the next element.
11       @return the traversed element
12      */
13      Object next();
14
15      /**
16       Tests if there is an element after the iterator position.
17       @return true if there is an element after the iterator position
18      */
19      boolean hasNext();
20
21      /**
22       Adds an element before the iterator position
23       and moves the iterator past the inserted element.
24       @param element the element to add
25      */
26      void add(Object element);
27
28      /**
29       Removes the last traversed element. This method may
30       only be called after a call to the next() method.
31      */
32      void remove();
33 }
```

Static Classes

- Every object of an inner class has a reference to the outer class.
 - It can access the instance variables and methods of the outer class.
- If an inner class does not need to access the data of the outer class,
 - It does not need a reference.
 - Declare it static to save the cost of the reference.
- **Example: Declare the `Node` class of the `LinkedList` class as `static`:**

```
public class LinkedList
{
    . . .
    static class Node
    {
        . . .
    }
}
```

Implementing Array Lists

- An array list maintains a reference to an array of elements.
- The array is large enough to hold all elements in the collection.
- When the array gets full, it is replaced by a larger one.
- An array list has an instance field that stores the current number of elements.

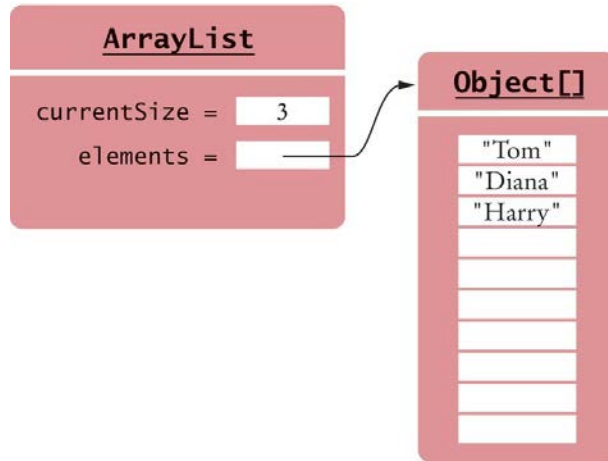


Figure 7 An Array List Stores its Elements in an Array

Implementing Array Lists

- Our ArrayList implementation will manage elements of type Object:

```
public class ArrayList
{
    private Object[] elements;
    private int currentSize;

    public ArrayList()
    {
        final int INITIAL_SIZE = 10;
        elements = new Object[INITIAL_SIZE];
        currentSize = 0;
    }

    public int size() { return currentSize; }
    . . .
}
```

Implementing Array Lists - Getting and Setting Elements

- Providing `get` and `set` methods:
 - Check for valid positions
 - Access the internal array at the given position
- Helper method to check bounds:

```
private void checkBounds(int n)
{
    if (n < 0 || n >= currentSize)
    {
        throw new IndexOutOfBoundsException();
    }
}
```


Implementing Array Lists - Getting and Setting Elements

- The `get` method:

```
public Object get(int pos)
{
    checkBounds(pos);
    return element[pos];
}
```

- The `set` method:

```
public void set(int pos, Object element)
{
    checkBounds(pos);
    elements[pos] = element;
}
```

- Getting and setting an element can be carried out with a bounded set of instructions, independent of the size of the array list.
- These are $O(1)$ operations.

Removing or Adding Elements

- To remove an element at position k , move the elements with higher index values.
- The remove method:

```
public Object remove(int pos)
{
    checkBounds(pos);
    Object removed = elements[pos];
    for (int i = pos + 1; i < currentSize; i++)
    {
        elements[i - 1] = elements[i];
    }
    currentSize--;
    return removed;
}
```

- On average, $n / 2$ elements need to move.
- Inserting a element also requires moving, on average, $n / 2$ elements.
- Inserting or removing an array list element is an $O(n)$ operation.

Removing or Adding Elements

- Exception: adding an element after the last element
 - Store the element in the array
 - Increment size
- An $O(1)$ operation
- A the `addLast` method:

```
public boolean addLast(Object newElement)
{
    growIfNecessary();
    currentSize++;
    elements[currentSize - 1] = newElement;
    return true;
}
```

Removing or Adding Elements

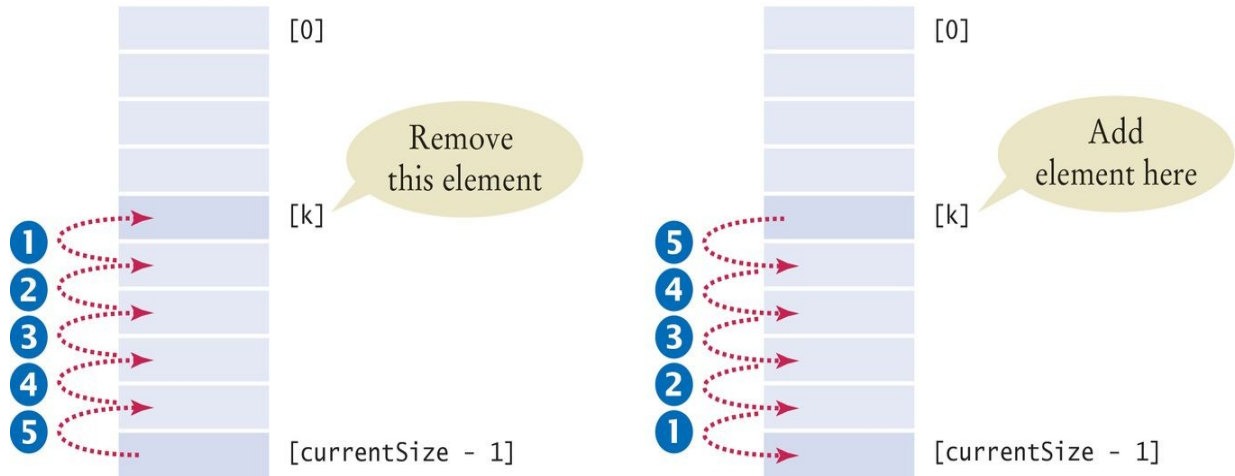


Figure 8 Removing and Adding Elements

Growing the Internal Array



© Craig Dingle/Stockphoto.

When an array list is completely full, we must move the contents to a larger array.

Growing the Internal Array

- When the array is full:
 - Create a bigger array
 - Copy the elements to the new array
 - New array replaces old
- Reallocation is $O(n)$.
- The `growIfNecessary` method:

```
private void growIfNecessary()
{
    if (currentSize == elements.length)
    {
        Object[] newElements = new Object[2 * elements.length]; ❶
        for (int i = 0; i < elements.length; i++)
        {
            newElements[i] = elements[i]; ❷
        }

        elements = newElements; ❸
    }
}
```

Growing the Internal Array

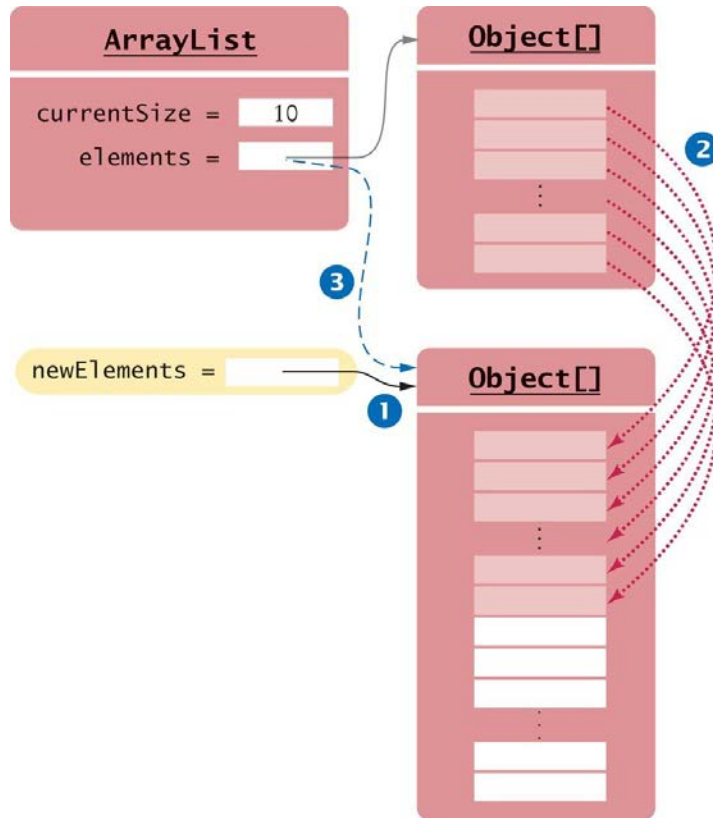


Figure 9 Reallocating the Internal Array

Growing the Internal Array

- Reallocation seldom happens.
- We amortize the cost of the reallocation over all the insertion or removals.
- Adding or removing the last element in an array list takes **amortized** $O(1)$ time. Written $O(1)+$

Efficiency of Array List and Linked List Operations

Table 2 Efficiency of Array List and Linked List Operations

Operation	Array List	Doubly-Linked List
Add/remove element at end.	$O(1)+$	$O(1)$
Add/remove element in the middle.	$O(n)$	$O(1)$
Get k th element.	$O(1)$	$O(k)$

Implementing Stacks and Queues


- Stacks and queues are abstract data types.
- We specify how operations must behave.
- We do not specify the implementation.
- Many different implementations are possible.

Stacks as Linked Lists

- A stack can be implemented as a sequence of nodes.
- New elements are “pushed” to one end of the sequence, and they are “popped” from the same end.
- Push and pop from the least expensive end - the front.
- The `push` and `pop` operations are identical to the `addFirst` and `removeFirst` operations of the linked list.

Stacks as Linked Lists

Adding an element



Removing an element




Figure 10 Push and Pop for a Stack Implemented as a Linked List

section_3_1/LinkedListStack.java

```
1  import java.util.NoSuchElementException;
2
3  /**
4   An implementation of a stack as a sequence of nodes.
5   */
6  public class LinkedListStack
7  {
8      private Node first;
9
10     /**
11      Constructs an empty stack.
12     */
13     public LinkedListStack()
14     {
15         first = null;
16     }
17
18     /**
19      Adds an element to the top of the stack.
20      @param element the element to add
21     */
22     public void push(Object element)
23     {
24         Node newNode = new Node();
25         newNode.data = element;
26         newNode.next = first;
27         first = newNode;
28     }
29
30     /**
31      Removes the element from the top of the stack.
32      @return the removed element
33     */
34     public Object pop()
35     {
```

Stacks as Arrays

- A stack can be implemented as an array.
- Push and pop from the least expensive end - the back.
The array must grow when it gets full.
- The `push` and `pop` operations are identical to the `addLast` and `removeLast` operations of an array list.
- `push` and `pop` are $O(1)$ operations.





Figure 11 A Stack Implemented as an Array

Queues as Linked Lists

- A queue can be implemented as a linked list:
 - Add elements at the back.
 - Remove elements at the front.
 - Keep a reference to last element.
- The `add` and `remove` operations are $O(1)$ operations.

Queues as Linked Lists

Adding an element



Removing an element




Figure 12 A Queue Implemented as a Linked List

Queues as Circular Arrays

- In a circular array, we wrap around to the beginning after the last element.

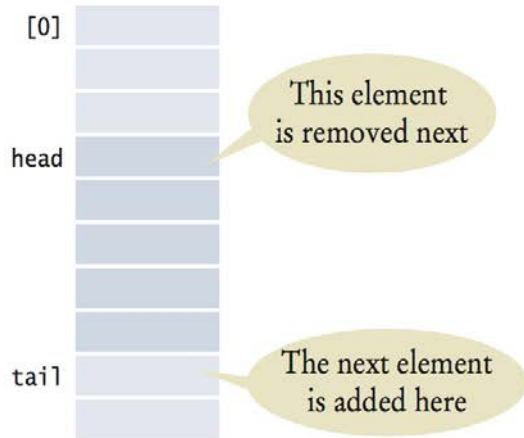


© ihsanyildizli/iStockphoto.

- When removing elements of a circular array,
 - increment the index at which the head of the queue is located.
- When the last element of the array is filled,
 - Wrap around and start storing at index 0
 - If elements have been removed there is room Else reallocate.
- All operations except reallocating are independent of the queue size
 - $O(1)$
- Reallocation is amortized constant time
 - $O(1)+$

Queues as Circular Arrays

Before wrapping around



After wrapping around




Figure 13 Queue Elements in a Circular Array

Queues as Circular Arrays

Table 3 Efficiency of Stack and Queue Operations

	Stack as Linked List	Stack as Array	Queue as Linked List	Queue as Circular Array
Add an element.	$O(1)$	$O(1)+$	$O(1)$	$O(1)+$
Remove an element.	$O(1)$	$O(1)+$	$O(1)$	$O(1)+$

section_3_4/CircularArrayQueue.java

```
1  import java.util.NoSuchElementException;
2
3  /**
4   An implementation of a queue as a circular array.
5   */
6  public class CircularArrayQueue
7  {
8      private Object[] elements;
9      private int currentSize;
10     private int head;
11     private int tail;
12
13     /**
14      Constructs an empty queue.
15     */
16     public CircularArrayQueue()
17     {
18         final int INITIAL_SIZE = 10;
19         elements = new
20             Object[INITIAL_SIZE];  currentSize
21             = 0;
22         head = 0;
23         tail = 0;
24     }
25
26     /**
27      Checks whether this queue is empty.
28      @return true if this queue is empty
29     */
30     public boolean empty() { return currentSize == 0; }
31
32     /**
33      Adds an element to the tail of this queue.
34      @param newElement the element to add
35     */
36     public void add(Object newElement)
```

Implementing a Hash Table

- In the Java library sets are implemented as hash sets and tree sets.
- **Hashing:** place items into an array at an index determined from the element.
- **Hash code:** an integer value that is computed from an object,
 - in such a way that different objects are likely to yield different hash codes.
- Collision: when two or more distinct objects have the same hash code. A good hash function minimizes collisions.
- A hash table uses the hash code to determine where to store each element.

Implementing a Hash Table

Table 4 Sample Strings and Their Hash Codes

String	Hash Code	String	Hash Code
"Adam"	2035631	"Juliet"	-2065036585
"Eve"	70068	"Katherine"	2079199209
"Harry"	69496448	"Sue"	83491
"Jim"	74478	"Ugh"	84982
"Joe"	74656	"VII"	84982

Hash Tables

- **Hash table:** An array that stores the set elements.
- **Hash code:** used as an array index into a hash table.
- Simplistic implementation
 - Very large array
 - Each object at its hashcode location
 - Simple to locate an element
 - But not practical




Figure 14 A Simplistic Implementation of a Hash Table

Hash Tables - **Realistic** Implementation

- A reasonable size array.
- Use the remainder operator to calculate the position.

```
int h = x.hashCode();  
if (h < 0) { h = -h; }  
position = h % arrayLength;
```

- Use **separate chaining** to handle collisions:
 - All colliding elements are collected in a linked list of elements with the same position value.
 - The lists are called **buckets**.
- Each entry of the hash table points to a sequence of nodes containing elements with the same hash code.
- A hash table can be implemented as an array of buckets—sequences of nodes that hold elements with the same hash code.

Hash Tables - **Realistic** Implementation

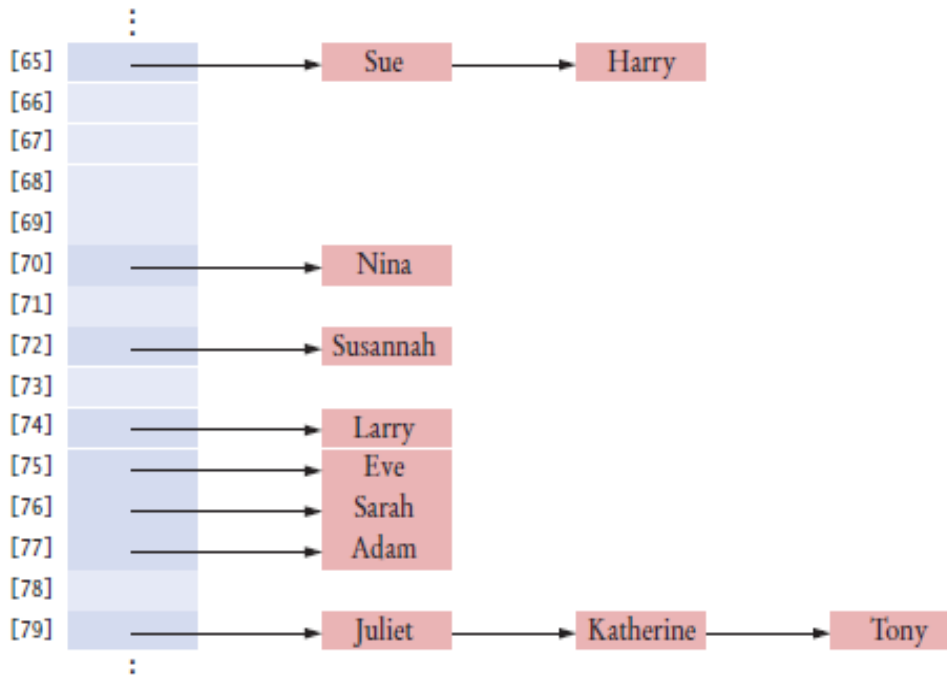


Figure 15 A Hash Table with Buckets to Store Elements with the Same Hash Code

Hash Tables



© Neil Kurtzman/iStockphoto.

Elements with the same hash code are placed in the same bucket.

Implementing a Hash Table - Finding an Element

- Algorithm to find an element, `obj`
 - Compute the hash code and compress it.
 - Gives an index `h` into the hash table.
 - Iterate through the elements of the bucket at position `h`.
 - Check element is equal to `obj`.
 - If a match is found among the elements of that bucket,
 - `obj` is in the set. Otherwise, it is not.
- If there are no or only a few collision:
 - adding, locating, and removing hash table elements takes $O(1)$ time.

Adding and Removing Elements

- Algorithm to add an element:
 - Compute the compressed hash code h .
 - Iterate through the elements of the bucket at position h .
 - For each element of the bucket, check whether it is equal to obj .
 - If a match is found among the elements of that bucket, then exit.
 - Otherwise, add a node containing obj to the beginning of the node sequence. If the load factor exceeds a fixed threshold, reallocate the table.
- Load factor: a measure of how full the table is.
 - The number of elements in the table divided by the table length.
- Adding an element to a hash table is $O(1)+$

Adding and Removing Elements

- Algorithm to remove an element:
 - Compute the hash code to find the bucket that should contain the object.
 - Try to find the element.
 - If it is present:
 - remove it.
 - otherwise, do nothing.
 - Shrink the table if it becomes too sparse.
- Removing an element from a hash table is $O(1)$ +

Iterating over a Hash Table

- When iterator points to the middle of a node chain,
 - easy to get the next element.
- When the iterator is at the end of a node chain,
 - Skip over empty buckets.
 - Advance the iterator to the first node of the first non-empty bucket.
- Iterator needs to store the bucket number and a reference to the current node in the node chain.

```
if (current != null && current.next != null) {
    current = current.next; // Move to next element in bucket
}
else // Move to next bucket
{
    do
    {
        bucketIndex++;
        if (bucketIndex == buckets.length)
            { throw new NoSuchElementException(); }
        current = buckets[bucketIndex];
    }
    while (current == null);
}
```

Iterating over a Hash Table

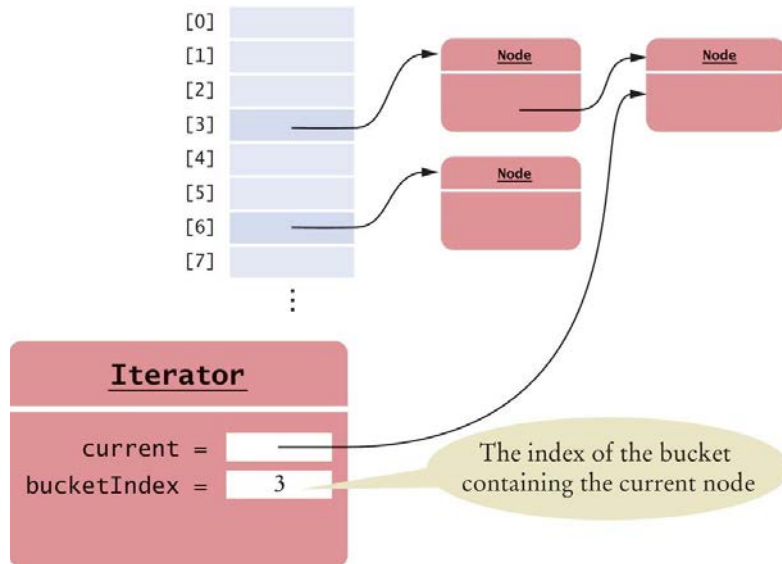


Figure 16 An Iterator to a Hash Table

Hash Table Efficiency

- The cost of iterating over all elements of a hash table:
 - Is proportional to the table length
 - Not the number of elements in the table
- Shrink the table when the load factor gets too small.
One iteration is $O(1)$.
- Iterating over the entire table is $O(n)$.

Table 5 Hash Table Efficiency	
Operation	Hash Table
Find an element.	$O(1)$
Add/remove an element.	$O(1)+$
Iterate through all elements.	$O(n)$

section_4/HashSet.java

```
1  import java.util.Iterator;
2  import java.util.NoSuchElementException;
3
4  /**
5   * This class implements a hash set using separate chaining.
6   */
7  public class HashSet
8  {
9      private Node[] buckets;
10     private int currentSize;
11
12     /**
13      * Constructs a hash table.
14      * @param bucketsLength the length of the buckets array
15      */
16     public HashSet(int bucketsLength)
17     {
18         buckets = new Node[bucketsLength];
19         currentSize = 0;
20     }
21
22     /**
23      * Tests for set membership.
24      * @param x an object
25      * @return true if x is an element of this set
26      */
27     public boolean contains(Object x)
28     {
29         int h = x.hashCode();
30         if (h < 0) { h = -h; }
31         h = h % buckets.length;
32
33         Node current = buckets[h];
34         while (current != null)
35         {
```

section_4/HashSetDemo.java

```
1  import java.util.Iterator;
2
3  /**
4   This program demonstrates the hash set class.
5   */
6  public class HashSetDemo
7  {
8      public static void main(String[] args)
9      {
10         HashSet names = new HashSet(101);
11
12         names.add("Harry");
13         names.add("Sue");
14         names.add("Nina");
15         names.add("Susannah");
16         ;
17         names.add("Larry");
18         names.add("Eve");
19         names.add("Sarah");
20         names.add("Adam");
21         names.add("Tony");
22         names.add("Katherine");
23         ;
24         names.add("Juliet");
25         names.add("Romeo");
26         names.remove("Romeo");
27         names.remove("George");
28         Iterator iter = names.iterator();
29         while (iter.hasNext())
30         {
31             System.out.println(iter.next());
32             ;
33         } } }
```

Program Run:

```
Harry  
Sue  
Nina  
Susannah  
Larry  
Eve  
Sarah  
Adam  
Juliet  
Katherine  
Tony
```